


# Introduction to R

## for I-TIPP Fellows

Matthew J Maenner, PhD  
13 October 2015

[www.github.com/mjmaenner/rdemo](http://www.github.com/mjmaenner/rdemo)




**"Using R is a bit akin to smoking. The beginning is difficult, one may get headaches and even gag the first few times. But in the long run, it becomes pleasurable and even addictive.**

**Yet, deep down, for those willing to be honest, there is something not fully healthy in it."**

**-Francois Pinard**

Text: <https://stat.ethz.ch/pipermail/r-help/2007-August/139047.html>

image: <http://electroniccigarettebuy.wordpress.com/2013/01/23/>



```
plot(ecdf(rnorm(n=10000, mean=1, sd=0.2)), xlab="Time", ylab="Ability", lwd=3, main="R Learning Curve", cex.lab=18/12)
```

# A very brief orientation to these things:

1. Interacting with R and RStudio
2. Basic syntax, notation, and operators
3. Doing things (functions), using packages
4. Strategies for learning and suggested tools to learn next

slides and code examples:

[www.github.com/mjmaenner/rdemo](https://www.github.com/mjmaenner/rdemo)

(feel free to interrupt and ask questions)

# R

Open-source programming language and statistical computing environment

Created in 1990's; implementation of the S language (Bell Labs)

The R Core Team now maintains "base" R, and 7000+ packages have been contributed by people all over the world.

Adapted from: [http://en.wikipedia.org/wiki/R\\_\(programming\\_language\)](http://en.wikipedia.org/wiki/R_(programming_language))

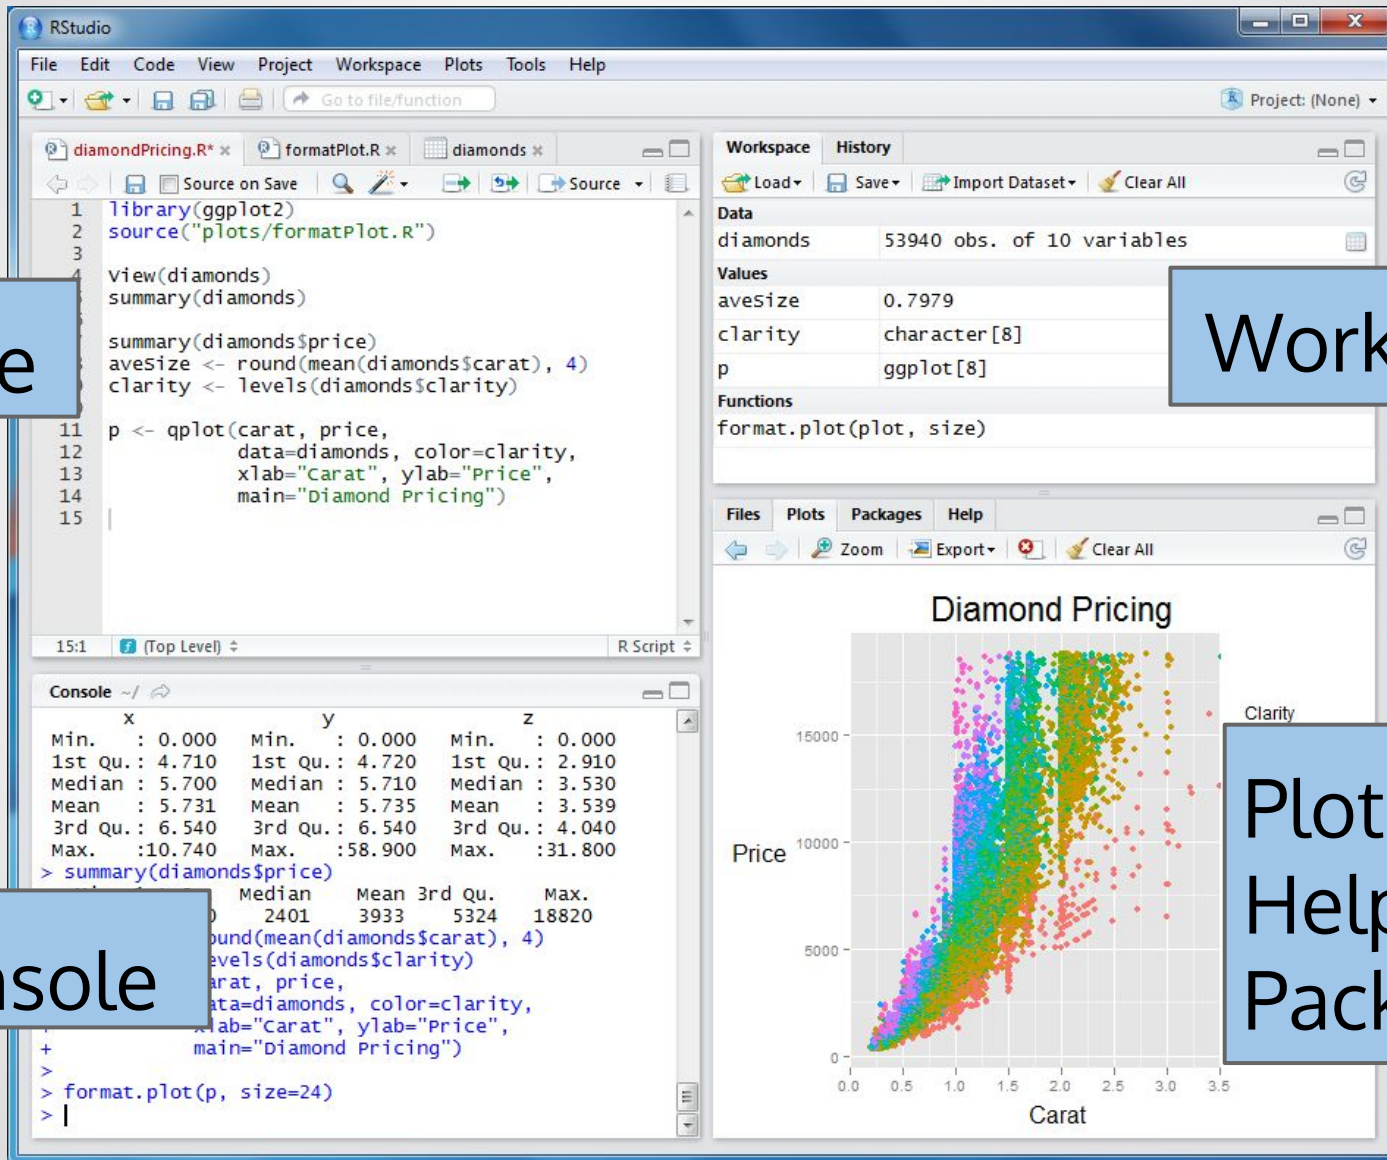
# RStudio

Source

Workspace

Console

Plots,  
Help,  
Packages



Default Rstudio setup <http://www.rstudio.com/ide/screenshots/>

# RStudio Basics

**Usually, type in the source window**

file -> new file -> R Script (if you don't see source window)

**Ctrl+Enter** submits line or selection to console

you can type in the console, too. Statements end at the end of lines, unless in a parentheses, bracket, or extended from previous line (like a + in ggplot2).

**Tab** completion for additional details:

type **summa** +tab to find **summary**; type **mean(** +tab to see options

At the R console, can recall previous commands with up arrow



# A simplification

**(most) everything is an object**  
either data or function

**Data objects store information.**

You do things *to* data objects

**Functions do things;** usually kept in packages

Call functions with **parentheses**: `citation()`  
without parentheses, it prints the function code

# Comprehensive R Archive Network

Perhaps the most compelling reason use R  
CRAN is like R's App Store

```
> install.packages("ggplot2") #from CRAN
Installing package(s) into 'C:/prog/R/R-3.2.1/library' ...
> library(ggplot2) #loads for use
```

You only have install 1st time (or to update package).  
Otherwise, only need `library()`

## What package(s) should I use?

- Find out what others use (Google, scientific literature)
- CRAN task views: <http://cran.r-project.org/web/views/>
- Serious packages often have their own papers.

# Basic Syntax and Operators

**#** comment (to end of line)

**<-** assign value to *objects*

inside functions, must assign *arguments* with "="

`weighted.mean(x=rate, w=population)`

**==** test of equality

**\$** specifies variables (or sub-objects) in a data.frame or list

**c** concatenates multiple values into a vector

`c(1, 2, 3)`

# How to get help

## ?function or help(function)

try it: ?mean, ?lm, ?summary, ?ggplot

it only works for packages that are currently loaded

## Google

specific task (e.g., calculate quantiles in R)

errors (e.g., "subscript out of bounds" R)

## More targeted searches:

?? performs a text search across packages that you have installed

try it: ??mean

RSiteSearch("ggplot2") searches R site for string

Find out other arguments and options: ?RSiteSearch

# Missing Values

Represented as **NA**; you should explicitly account for NA values when transforming or subsetting data.

Try this: `sex<- c(1, 0, NA)`

```
> sex==1
```

```
[1] TRUE FALSE NA #NA does not evaluate
```

also, notice the vector recycling

Try: `is.na(sex)` or `!(is.na(sex))` #What happens?  
`is.finite(sex)`

more: <http://www.statmethods.net/input/missingdata.html>

# Loading Data

Many people prefer **CSV**

```
crimeat1 <- read.csv("~/crime.csv", header=TRUE)
```

*read.csv is a special case of read.table*

**SPSS, Stata** and **SAS** via the **foreign** & **sas7bdat** packages

```
data<- read.spss(file="c:/data.sav",to.data.frame=TRUE)
```

**XLS or XLSX** (not recommended, via **xlsx** package)

```
data<-read.xlsx(file="c:/data.xlsx",sheetIndex=1)
```

Database connectivity via **RODBC** and others.

```
ch<-odbcConnect(dsn="SQLServer11", readOnly=TRUE)  
data<-sqlFetch(channel=ch, sqtable='_DR_FINALDATA')
```

# Loading Data (exercise)

local csv:

```
> crimeat1 <- read.csv("~/crime.csv",  
header=TRUE)
```

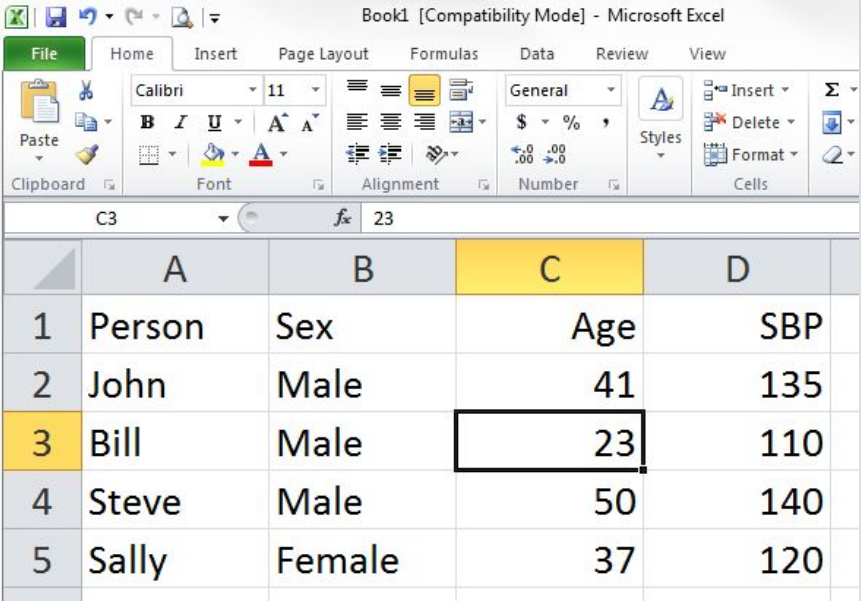
**RSTUDIO GUI:** environment tab (upper right)  
click "import dataset" -- "from text file"  
follow prompts

# The data.frame

A rectangular data set  
often abbreviated as "df"

**Row** : an observation

**Column** : a variable



	A	B	C	D
1	Person	Sex	Age	SBP
2	John	Male	41	135
3	Bill	Male	23	110
4	Steve	Male	50	140
5	Sally	Female	37	120

each column has its own **class**:

- **numeric** values
- **character** strings
- **factors** (categorical; may or may not be ordered)

Check the class of your variable:

```
> class(crimeat1$beat)
```

```
> [1] "factor"
```



# Look at your data

Did it load the way you thought it would?

```
> dim(crimeatl)
```

```
[1] 230765 25 # number of rows and columns
```

```
> str(crimeatl)
```

```
'data.frame': 230765 obs. of 25 variables:
 $ MI_PRINX      : Factor w/ 230779 levels "-84.36061", "-84.37037",...: 7 8 9 10 11 12 13 14 15 16 ...
 $ offense_id    : num  90360664 90370891 91681984 72692336 80081069 ...
 $ rpt_date      : Factor w/ 2474 levels "", "01/01/2009",...: 247 254 1172 381 1950 401 93
142 254 451 ...
```

# Look at your data

Did it load the way you thought it would?

```
> View(crimeatl) # opens spreadsheet like view
```

```
> colnames(crimeatl)
```

[1]	"MI_PRINX"	"offense_id"	"rpt_date"	"occur_date"
[5]	"occur_time"	"poss_date"	"poss_time"	"beat"
[9]	"apt_office_prefix"	"apt_office_num"	"location"	"MinOfucr"
[13]	"MinOfibr_code"	"dispo_code"	"MaxOfnum_victims"	"Shift"
[17]	"Avg.Day"	"loc_type"	"UC2.Literal"	"neighborhood"
[21]	"npu"	"x"	"y"	

# Look at your data (cont'd)

```
> colnames(crimeatl)
[1] "MI_PRINX"      "offense_id"      "rpt_date"        "occur_date"
[5] "occur_time"     "poss_date"       "poss_time"       "beat"
[9] "apt_office_prefix" "apt_office_num"  "location"        "MinOfucr"
[13] "MinOfibr_code"   "dispo_code"      "MaxOfnum_victims" "Shift"
[17] "Avg.Day"         "loc_type"        "UC2.Literal"     "neighborhood"
[21] "npu"            "x"               "y"
```

```
> summary(crimeatl)
```

MI_PRINX		offense_id		rpt_date	
1160569:	1	Min. :	7.269e+07	11/17/2009	: 171
1160570:	1	1st Qu.:	1.020e+08	06/01/2009	: 154
1160572:	1	Median :	1.206e+08	07/14/2009	: 154
1160573:	1	Mean :	4.249e+08	08/29/2011	: 154
1160574:	1	3rd Qu.:	1.333e+08	07/27/2009	: 152
(Other):	230733	Max. :	1.527e+11	(Other)	: 229953
NA's :	27	NA's :	27	NA's	: 27

... <truncated>

```
> head(crimeatl)
```

	MI_PRINX	offense_id		rpt_date		occur_date
1	1160569	90360664	02/05/2009		02/03/2009	
2	1160570	90370891	02/06/2009		02/06/2009	
3	1160572	91681984	06/17/2009		06/17/2009	
4	1160573	72692336	02/24/2010		02/24/2010	
5	1160574	80081069	10/06/2010		01/08/2008	
6	1160575	82040835	02/27/2009		07/21/2008	

# Using Functions

```
function(arg1 = value, arg2 = value, ...)
```

- Follow the arguments (help or tab complete)
- Explicitly name arguments, especially when using something new
  - R will also do partial matching
- Decide where output should go:
  - console, or to be saved to a new object

# Functions

```
function(arg1 = value, arg2 = value, ...)
```

```
day <- table(crimeat1$Avg.Day)
```

# Functions

```
function(arg1 = value, arg2 = value, ...)
```

```
day <- table(crimeat1$Avg.Day)
```



**function**

# Functions

```
function(arg1 = value, arg2 = value, ...)
```

**Data as argument**

```
day <- table(crimeat1$Avg.Day)
```

**function**

# Functions

```
function(arg1 = value, arg2 = value, ...)
```

**Data as argument**

```
day <- table(crimeat1$Avg.Day)
```

**function**

**assign object to store output  
(otherwise it prints default output)**



# Useful Descriptive Functions

```
> table(crimeat1$Avg.Day)
```

	1	Day	Eve	Fri	Mon	Morn	NULL	Sat	Sun	Thu	Tue	Unk	Wed
0	0	0	0	30728	29668	0	0	31292	27701	29653	30565	21361	29770

#try adding ,useNA= “always”

**aggregate()** # performing function on groups

**mean()**

**quantile()**

# Accessing parts of a data frame

`data.frame[row, column]`

`crimeat1[ 1:3 , ]` # returns rows 1-3

`crimeat1[ , 1:3 ]` # returns columns 1-3

These return the same values:

`crimeat1$Avg.Day` #tab-complete after \$

`crimeat1[ , 17 ]`

`crimeat1[ , 'Avg.Day' ]`

`crimeat1[ 'Avg.Day' ]` # !!! returns df!

# example

We are going to make a map of reported crimes, so I know where to buy a house in Atlanta.

```
> summary(crimeatl[,c("x","y")])
```

x	y
-84.36212: 2945	33.84676: 2945
-84.41278: 1768	33.80388: 1751
-84.40902: 838	33.68274: 856
-84.43276: 835	33.73812: 842
-84.49773: 831	33.68677: 834
(Other) :223521	(Other) :223510
NA's : 27	NA's : 27

```
> class(crimeatl$x)
```

```
[1] "factor"
```

# example (cont'd)

We need to convert the factor variables to numeric:

```
crimeat1$lon<-as.numeric(as.character(crimeat1$x))
```

```
crimeat1$lat<-as.numeric(as.character(crimeat1$y))
```

Hint: always go factor -> character -> numeric

# example (cont'd)

```
> summary(crimeatl[, c("lon", "lat")])
```

lon		lat	
Min.	:-84.55	Min.	:-84.50
1st Qu.	:-84.43	1st Qu.	: 33.73
Median	:-84.40	Median	: 33.76
Mean	:-84.41	Mean	: 33.75
3rd Qu.	:-84.37	3rd Qu.	: 33.78
Max.	:-84.29	Max.	: 33.89
NA's	:41	NA's	:27

Dropping the outlier lat value(s)

```
crimeatl<- crimeatl[ crimeatl$lat > 33, ]
```

# example (cont'd)

In order to make a map, we're going to need some functions that are stored in packages.

If this is your first time, run

```
install.packages(c("ggplot2", "ggmap", "ggthemes", "lubridate"))
```

to download everything.

Then, load the packages with `library()`

```
library(ggplot2)
```


```
library(scales)
```

```
library(ggmap)
```

```
library(ggthemes)
```


# example (cont'd)

```
ggplot(data=crimeatl, aes(x=lon, y=lat)) +  
  geom_point()
```



# example (cont'd)

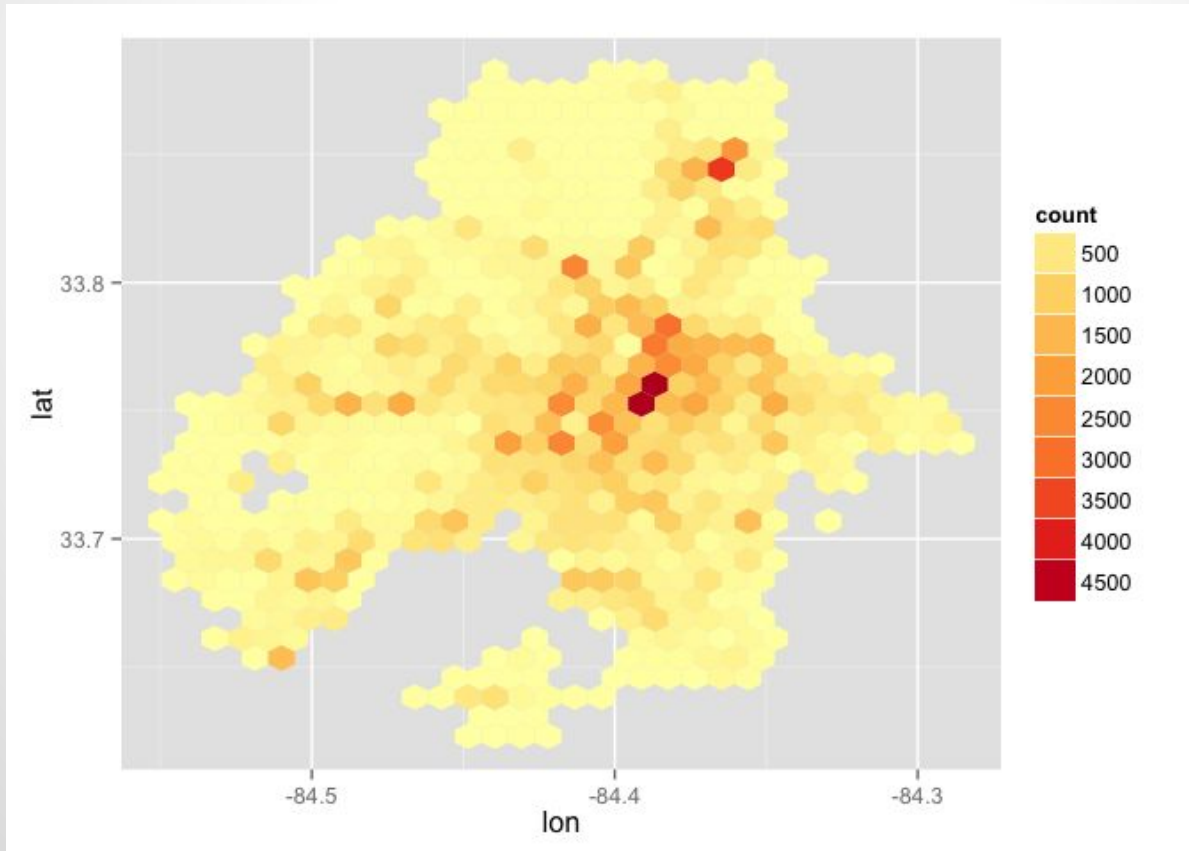
```
ggplot(data=crimeatl, aes(x=lon, y=lat)) +  
  geom_hex()
```





# example (cont'd)

```
ggplot(data=crimeatl, aes(x=lon, y=lat)) +  
  geom_hex() +  
  scale_fill_distiller(palette="YlOrRd",  
    breaks=pretty_breaks(n=10))
```



# example (cont'd)

We need a map to go with the data. First, we will get the bounds of the coordinates we need:

```
> r.lon <- range(crimeatl$lon, na.rm=TRUE)
> r.lat <- range(crimeatl$lat, na.rm=TRUE)
> bounds<-c(r.lon[1], r.lat[1], r.lon[2], r.lat[2])
> bounds
[1] -84.55049  33.62172 -84.28641  33.88613
```

Using the bounds, we request a map:

```
atl.map <- get_map(location=bounds, maptype = "toner",
                   zoom=13, crop=FALSE)
```

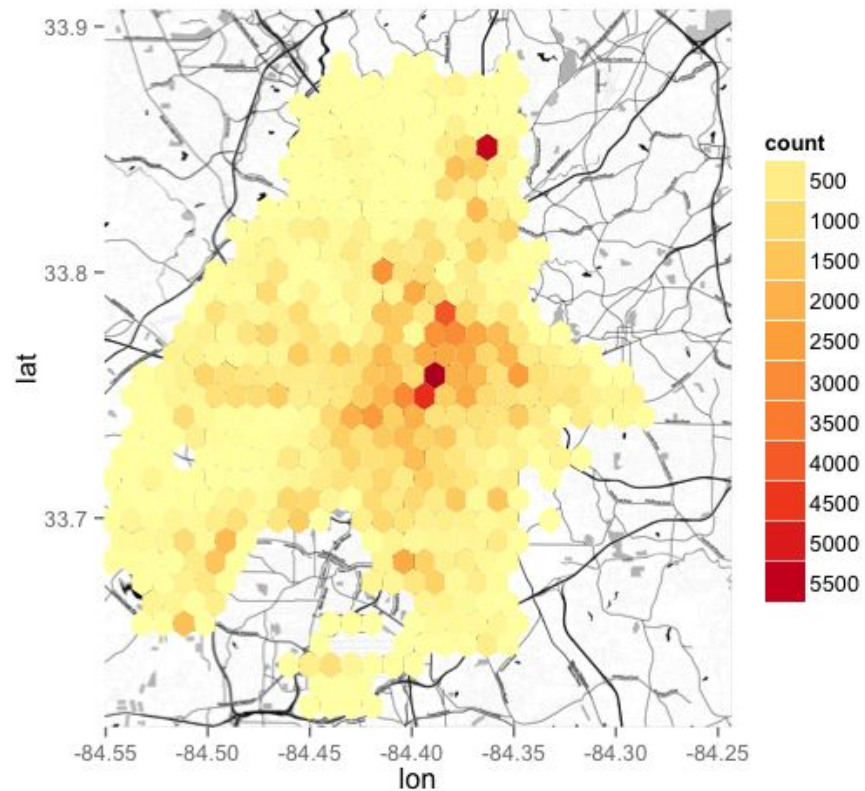
# example (cont'd)

```
ggmap(atl.map)
```



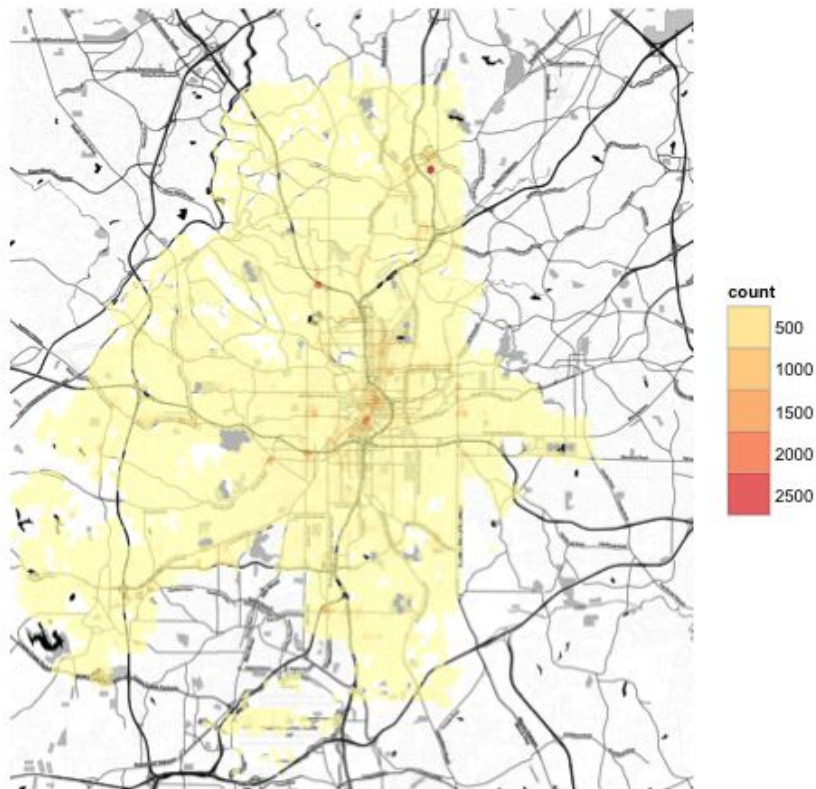
# example (cont'd)

```
ggmap(atl.map) +  
  geom_hex(data=crimeatl, aes(x=lon, y=lat))+  
  scale_fill_distiller(palette="YlOrRd",  
                      breaks=pretty_breaks(n=10))+
```



# example (cont'd)

```
ggmap(atl.map) +  
  geom_hex(data=crimeatl, aes(x=lon, y=lat), alpha=.7, bins=100) +  
  scale_fill_distiller(palette="YlOrRd",  
                      breaks=pretty_breaks(n=10)) +  
  theme_map()+theme(legend.position="right")
```



# example (cont'd)

Suppose we're mostly interested in homicides:

```
murderat1 <- crimeat1[ crimeat1$UC2.Literal == "HOMICIDE",]
```

and we want to know when they happen

```
> library(lubridate)
> murderat1$murderdate <- mdy(murderat1$occur_date)
> murderat1$murderyear <- year(murderat1$murderdate)
>
> table(murderat1$murderyear)
```

2001	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
1	2	2	2	76	88	88	82	83	92	69



# example (cont'd)

69 murders so far this year... through week 40

```
> table(murderat1$murderyear)
```

2001	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
1	2	2	2	76	88	88	82	83	92	69

```
> 69/(40/52)
```

```
[1] 89.7
```

# example (cont'd)

Can also look by month

```
> murderatl$murdermonth<-month(murderatl$murderdate)  
> table(murderatl$murderyear, murderatl$murdermonth)
```

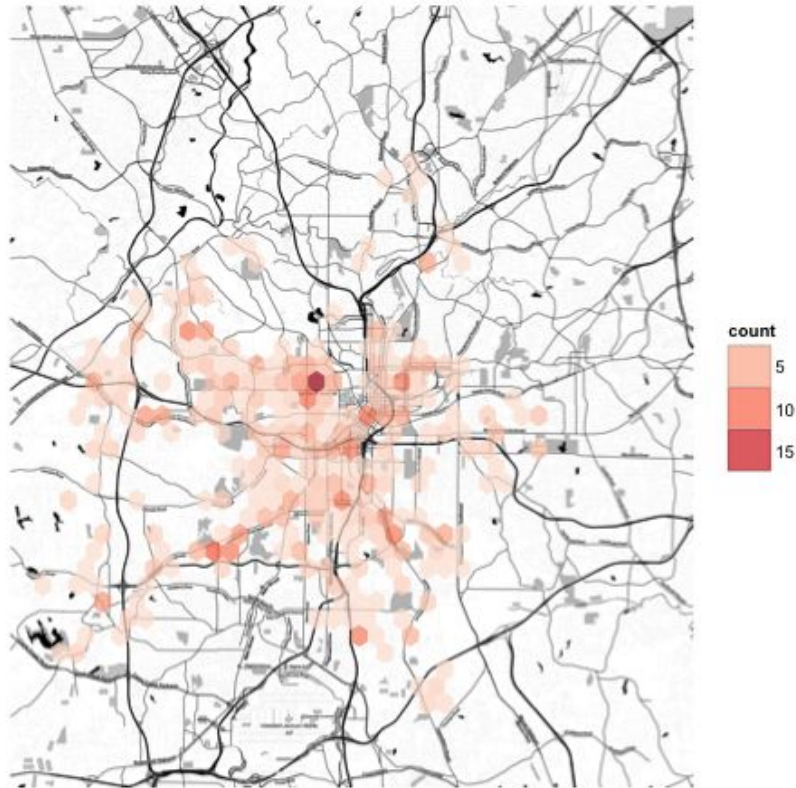
	1	2	3	4	5	6	7	8	9	10	11	12
2001	0	0	0	0	0	0	0	1	0	0	0	0
2006	0	0	0	1	0	0	1	0	0	0	0	0
2007	0	1	0	0	0	0	0	1	0	0	0	0
2008	0	0	0	1	0	0	0	0	1	0	0	0
2009	6	9	5	5	9	9	7	6	4	8	1	7
2010	5	4	9	9	6	6	7	10	11	9	6	6
2011	6	3	6	7	8	11	7	4	8	7	10	11
2012	5	3	10	5	5	9	8	8	6	10	8	5
2013	7	7	3	6	7	7	8	9	5	7	10	7
2014	6	1	12	4	4	12	8	8	10	9	12	6
2015	6	8	5	8	11	5	14	4	7	1	0	0



# example (cont'd)

Map of homicides only:

```
ggmap(atl.map) +  
  geom_hex(data=murderatl, aes(x=lon, y=lat), alpha=.7, bins=40) +  
  scale_fill_distiller(palette="Reds")+  
  theme_map()+theme(legend.position="right")
```



# Where to go from here?

- **Practice, experiment, make mistakes**
- find packages that make your work easier
  - tutorials ("vignettes") often available
- **?help or help(function)**
- google your errors
  - stack exchange
  - r mailing list

# The “Hadleyverse”

Hadley and his colleagues make some of the most popular packages. Some serve as elegant wrappers for existing R functionality, others make calls to functions written in C.

**reshape2 & tidyR** - reshaping data (e.g., from 'wide' to 'long')

**plyr & dplyr** - transform data using split-apply-combine framework

**lubridate** - magically work with dates, times, durations

**stringr** - manipulate character strings

**ggplot2** - amazing graphing packages

**rvest, devtools, haven... lots**

# Excellent resources to help along the way...

Jenny Bryan's STAT545 course materials

<http://www.stat.ubc.ca/~jenny/STAT545A/current.html>

R reference card, <http://cran.r-project.org/doc/contrib/Short-refcard.pdf> PDF cheatsheet

**R Cheat Sheets:** <https://drive.google.com/folderview?pli=1&id=0ByIrlIAE4KMTtcVBmdm1BOEZoeEk#>

Quick-R, <http://www.statmethods.net/> : Excellent Info on basic functions

**R Twotutorials**, <http://www.twotutorials.com/>: Two-minute "how-to" presentations

**R Cookbook**, <http://wiki.stdout.org/rcookbook/>: Great R info in general, especially graphics

R for SPSS & SAS users, <https://sites.google.com/site/r4statistics/books/free-version>:

PDF for transitioning to R

[https://science.nature.nps.gov/im/datamgmt/statistics/R/documents/R\\_for\\_SAS\\_SPSS\\_users.pdf](https://science.nature.nps.gov/im/datamgmt/statistics/R/documents/R_for_SAS_SPSS_users.pdf)

**swirl package** (on CRAN and [www.swirlstats.com](http://www.swirlstats.com))

R Inferno: The 9 circles of R Hell [http://www.burns-stat.com/pages/Tutor/R\\_inferno.pdf](http://www.burns-stat.com/pages/Tutor/R_inferno.pdf)

slides and code examples:

[www.github.com/mjmaenner/rdemo](https://www.github.com/mjmaenner/rdemo)